

# PIPELINES WITH SLURM

Bio Code club end-of-semester  
event

# BASH AND SLURM

Conventional script:

```
> bash scriptname.bash
```

script is executed

On HPC:

```
> sbatch scriptname.bash
```

1)resources requested, job queued

2)job allocated on compute node

3)script is executed



Compute cluster

# SCALING UP ON HPC

## Benefits

- More time
- More memory
- Larger/more datasets

## Requirements

- Say you want to scale from 10 input files to 50 or 100
- But you have a complex pipelines (branching, multiple steps)
- Learning Workflow tools like Nextflow are a whole new system
- Basic → intermediate level scripting might be enough to do the coordination

## This talk

- Introduce some features in SLURM to parallelize and coordinate scripts
- Talk about some debugging techniques
- Tools and Templates

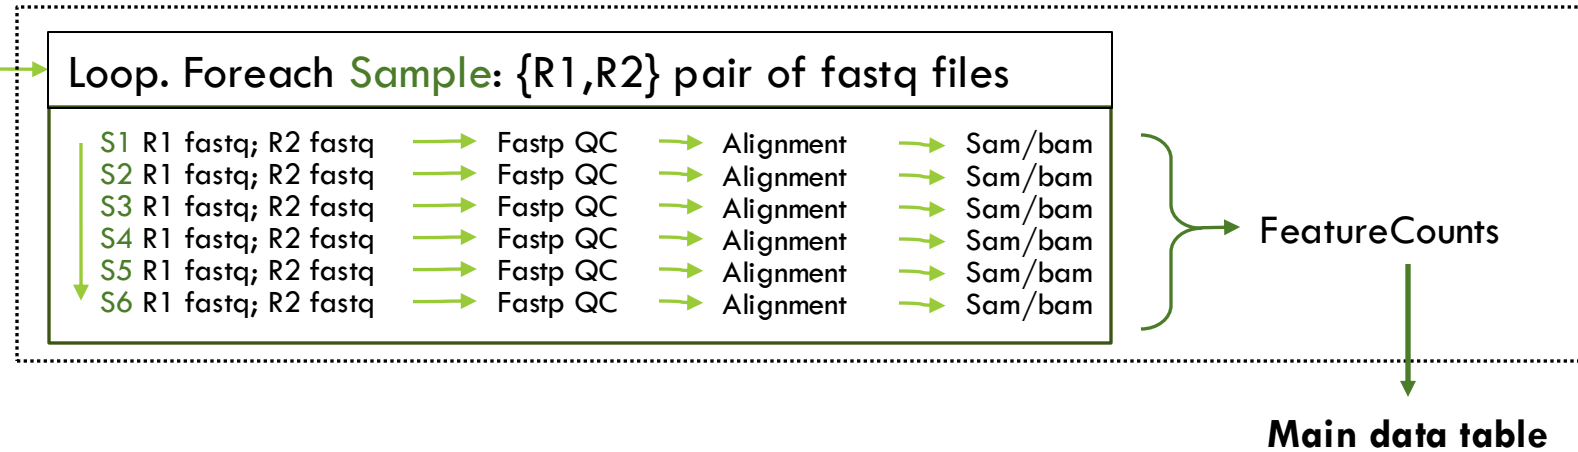


Compute cluster

# All data processed in single script, single job

## Single RNA-seq Analysis script

RNA-seq fastq data



## Pipeline synopsis

- Several steps per input file set **S1...S6**
- Loop iteration contains bulk of work, but ...
- ...not parallelized: doesn't scale well to large/many files
- Final step FeatureCounts run after loop has completed

# All data processed in single script, single job

## Loop. Sequential execution

S1	R1 fastq; R2 fastq	→ 20mi Fastp QC	→ 36min Alignment	→ 16min Sam/bam	72min
S2	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
S3	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
S4	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
S5	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
S6	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
					<hr/> 7 hours, 12 minutes

## Loop. 2x cores

S1	→ 10mi Fastp QC	→ 12min Alignment	→ 8min Sam/bam	36min
S2	→ Fastp QC	→ Alignment	→ Sam/bam	36min
S3	→ Fastp QC	→ Alignment	→ Sam/bam	36min
S4	→ Fastp QC	→ Alignment	→ Sam/bam	36min
S5	→ Fastp QC	→ Alignment	→ Sam/bam	36min
S6	→ Fastp QC	→ Alignment	→ Sam/bam	36min
				<hr/> 3 hours, 36 minutes

Cores:

Per sample time (minutes):

Total time (hours:minutes) for 6 iterations:

1x	2x	4x*	8x*
72	36	24	20
7:12	3:36	2:24	2:00

*\*diminishing return*

# ALL DATA PROCESSED SEQUENTIALLY

## Loop. Sequential execution

S1	R1 fastq; R2 fastq	→ 20min Fastp QC	→ 36min Alignment	→ 16min Sam/bam	72min
S2	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
S3	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
S4	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
S5	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
S6	R1 fastq; R2 fastq	→ Fastp QC	→ Alignment	→ Sam/bam	72min
					7 hours, 12 minutes

## Loop. 2x cores

S1	→ 10min Fastp QC	→ 12min Alignment	→ 8min Sam/bam	36min
S2	→ Fastp QC	→ Alignment	→ Sam/bam	36min
S3	→ Fastp QC	→ Alignment	→ Sam/bam	36min
S4	→ Fastp QC	→ Alignment	→ Sam/bam	36min
S5	→ Fastp QC	→ Alignment	→ Sam/bam	36min
S6	→ Fastp QC	→ Alignment	→ Sam/bam	36min
				3 hours, 36 minutes

Cores:

Per sample time (minutes):

Total time (hours:minutes) for 6 iterations:

1x	2x	4x*	8x*
72	36	24	20
7:12	3:36	2:24	2:00

*\*diminishing return*

# PARALLEL TIMING

				1x	2x	4x	8x
S1	R1 fastq; R2 fastq	→	Fastp QC → Alignment → Sam/bam	72min	36min	24min	20min
S2	R1 fastq; R2 fastq	→	Fastp QC → Alignment → Sam/bam	72min	36min	24min	20min
S3	R1 fastq; R2 fastq	→	Fastp QC → Alignment → Sam/bam	72min	36min	24min	20min
S4	R1 fastq; R2 fastq	→	Fastp QC → Alignment → Sam/bam	72min	36min	24min	20min
S5	R1 fastq; R2 fastq	→	Fastp QC → Alignment → Sam/bam	72min	36min	24min	20min
S6	R1 fastq; R2 fastq	→	Fastp QC → Alignment → Sam/bam	72min	36min	24min	20min
				<b>72min</b>	<b>36min</b>	<b>24min</b>	<b>20min</b>

- Put independent work in independent jobs:
  - Riviera: 40 Nodes, 7912 cores
  - Alpine: 696 Nodes, 40,200 cores
  - HPC!!!
- Parallelization
  - In a bash script: distribute jobs across cores
  - Scaling up:
    - Total run time is largest individual job, not sum of jobs
    - Dependent on queue- resources/priority

# Using sbatch with the array argument

```
$ sbatch --array=1-6 main.sh
```

<- hit enter

Submitted job with ID 9223918

Launches Main RNA-seq Analysis script *6 times*;  
sets `SLURM_ARRAY_TASK_ID` to different values: (1,2,3,4,5,6)

`SLURM_ARRAY_TASK_ID=1`

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

`SLURM_ARRAY_TASK_ID=2`

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

`SLURM_ARRAY_TASK_ID=3`

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

`SLURM_ARRAY_TASK_ID=4`

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

`SLURM_ARRAY_TASK_ID=5`

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

`SLURM_ARRAY_TASK_ID=6`

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
25736_1	short-cpu	rna-seq	dking	R	0:04	1	node001
25736_2	short-cpu	rna-seq	dking	R	0:04	1	node001
25736_3	short-cpu	rna-seq	dking	R	0:04	1	node001
25736_4	short-cpu	rna-seq	dking	R	0:01	1	node001
25736_5	short-cpu	rna-seq	dking	R	0:01	1	node001
25736_6	short-cpu	rna-seq	dking	R	0:01	1	node001



# Using sbatch with the array argument

```
$ sbatch --array=1-6 main.sh
```

<- hit enter

Submitted job with ID 9223918

SLURM\_ARRAY\_TASK\_ID=1

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

SLURM\_ARRAY\_TASK\_ID=2

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

SLURM\_ARRAY\_TASK\_ID=3

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

SLURM\_ARRAY\_TASK\_ID=4

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

SLURM\_ARRAY\_TASK\_ID=5

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

SLURM\_ARRAY\_TASK\_ID=6

S ■ R1 fastq; R2 fastq → Fastp QC → Alignment → Sam/bam

Launches Main RNA-seq Analysis script *6 times*;  
sets SLURM\_ARRAY\_TASK\_ID to different values: (1,2,3,4,5,6)

What's this? →

S ■ R1 fastq; R2 fastq

Base input filenames on SLURM\_ARRAY\_TASK\_ID

**e.g.:**

R1=S\${SLURM\_ARRAY\_TASK\_ID}\_R1.fastq

R2=S\${SLURM\_ARRAY\_TASK\_ID}\_R2.fastq

fastp -i \$R1 -l \$R2



SLURM\_ARRAY\_TASK\_ID=3

R1=S3\_R1.fastq

R2=S3\_R2.fastq

fastp -i S3\_R1.fastq -l S3\_R2.fastq

# Features of array jobs:

## monitoring with **squeue**

```
$ squeue -u $USER
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
25736_1	short-cpu	rna-seq	dking	R	0:04	1	node001
25736_2	short-cpu	rna-seq	dking	R	0:04	1	node001
25736_3	short-cpu	rna-seq	dking	R	0:04	1	node001
25736_4	short-cpu	rna-seq	dking	R	0:01	1	node001
25736_5	short-cpu	rna-seq	dking	R	0:01	1	node001
25736_6	short-cpu	rna-seq	dking	R	0:01	1	node001

Run as a group,  
Share JOB\_ID  
Number after '\_' is array job ID.

Might be staggered  
due to resource/priority

# Features of array jobs: Specifying log file names

## Single job script

single.sh

```
#!/usr/bin/env bash
#SBATCH ...
#SBATCH --job-name=rna-seq
#SBATCH --output=%x.%j.log
#SBATCH ...
...
```

## Log file:

rna-seq.25736.log  
%x %j

## Array job script

array.sh

```
#!/usr/bin/env bash
#SBATCH ...
#SBATCH --job-name=rna-seq
#SBATCH --output=%x.%A_%a.log
#SBATCH ...
...
```

rna-seq.25736\_1.log  
rna-seq.25736\_2.log  
rna-seq.25736\_3.log  
rna-seq.25736\_4.log  
rna-seq.25736\_5.log  
rna-seq.25736\_6.log  
%x %A %a

# Features of array jobs: filename patterns for logs

## FILENAME PATTERN

**sbatch** allows for a filename pattern to contain one or more replacement symbols, which are a percent sign "%" followed by a letter (e.g. %j).

**%A** Job array's master job allocation number.

**%a** Job array ID (index) number.

**%j** jobid of the running job.

**%x** Job name.

--From "sbatch" help page

rna-seq.25736.log

%x %j

rna-seq.25736\_1.log

rna-seq.25736\_2.log

rna-seq.25736\_3.log

rna-seq.25736\_4.log

rna-seq.25736\_5.log

rna-seq.25736\_6.log

%x %A %a

# SUBMISSION SCRIPT VERSUS ARRAY


submit.sh

```
#!/usr/bin/env bash
# run as `bash submit.sh`

for inputfile in *.R1.fastq
do
    sbatch single_rnaseq_analyzer.sh $inputfile
done
```

Simpler script, but no longer tied to a shared job id.

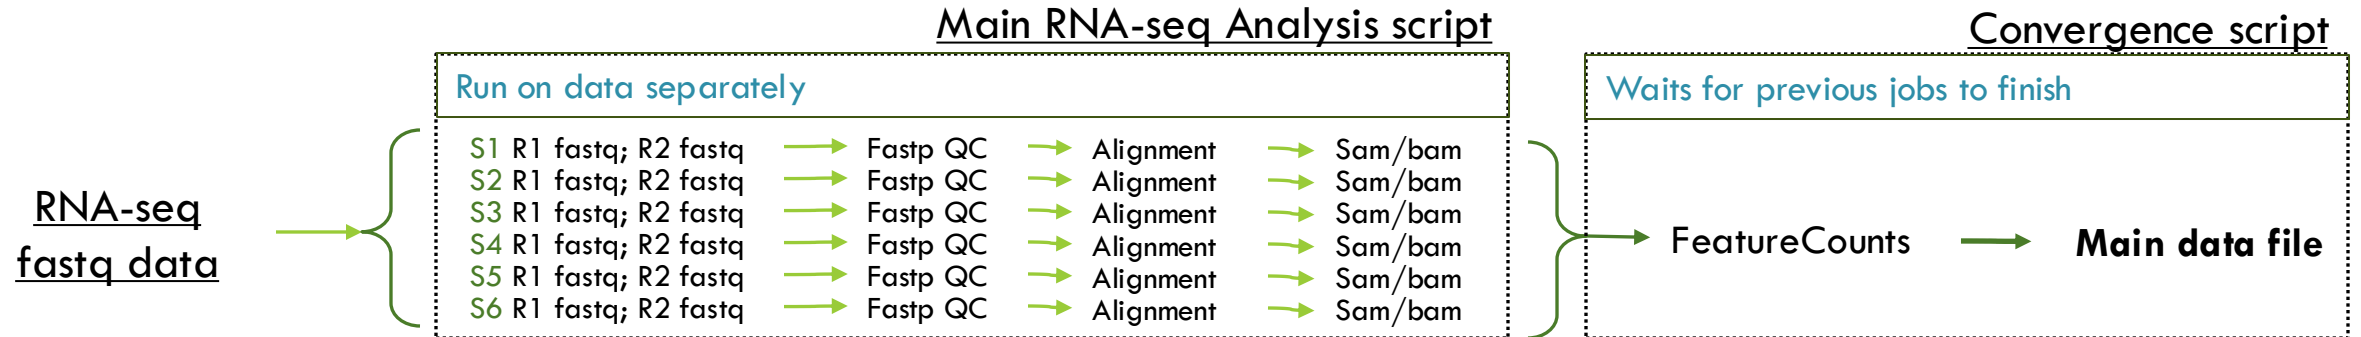
Harder to track, coordinate with other jobs



# JOB COORDINATION

Converge after array job

# Parallel workflow $\ni$ Convergence



## How?

```
$ sbatch --array=1-6 main.sh
```

Submitted job with ID 9223918

```
$ sbatch --dependency=afterok:9223918 converge.sh
```

Submitted job with ID 9223925

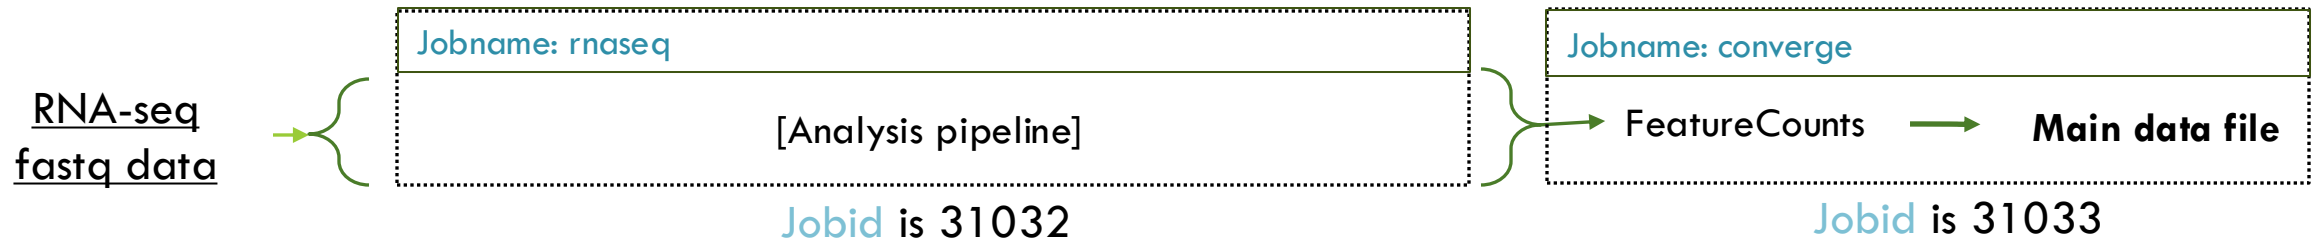
## Capture job id

```
$ jobid=$(sbatch -p --array=1-6 main.sh)
```

```
$ sbatch --dependency=afterok:$jobid converge.sh
```

Submitted job with ID 9223925

# SLURM DEPENDENCIES



```
$ jobid=$(sbatch -p --array=1-6 main.sh)
$ sbatch --dependency=afterok:$jobid converge.sh
Submitted job with ID 31033
```

```
$ squeue -u $USER
JOBID PARTITION  NAME   USER ST TIME  NODES NODELIST(REASON)
31033  short-cpu converge dking PD 0:00  1 (Dependency)
31032_1 short-cpu rnaseq   dking R 0:01  1 node001
31032_2 short-cpu rnaseq   dking R 0:01  1 node001
31032_3 short-cpu rnaseq   dking R 0:01  1 node001
```

...



# SLURM DEPENDENCIES

**-d, --dependency=<dependency\_list>**

- Defer the start of this job until the specified dependencies have been satisfied.
- <dependency\_list> is of the form

**afterok:job\_id** -Job with job\_id must complete successfully

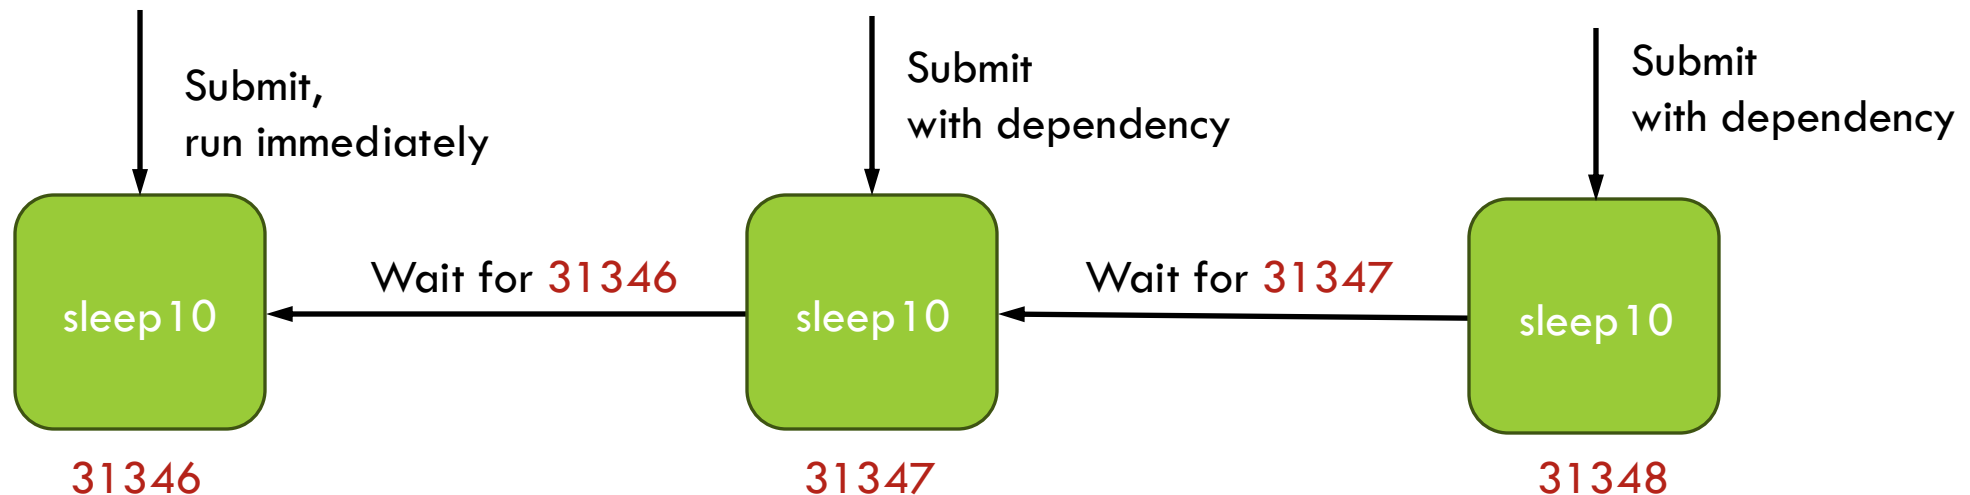
**afternotok:job\_id** -Job with job\_id failed: (non-zero exit code, node failure, timed out, etc). *Cleanup/error script?*

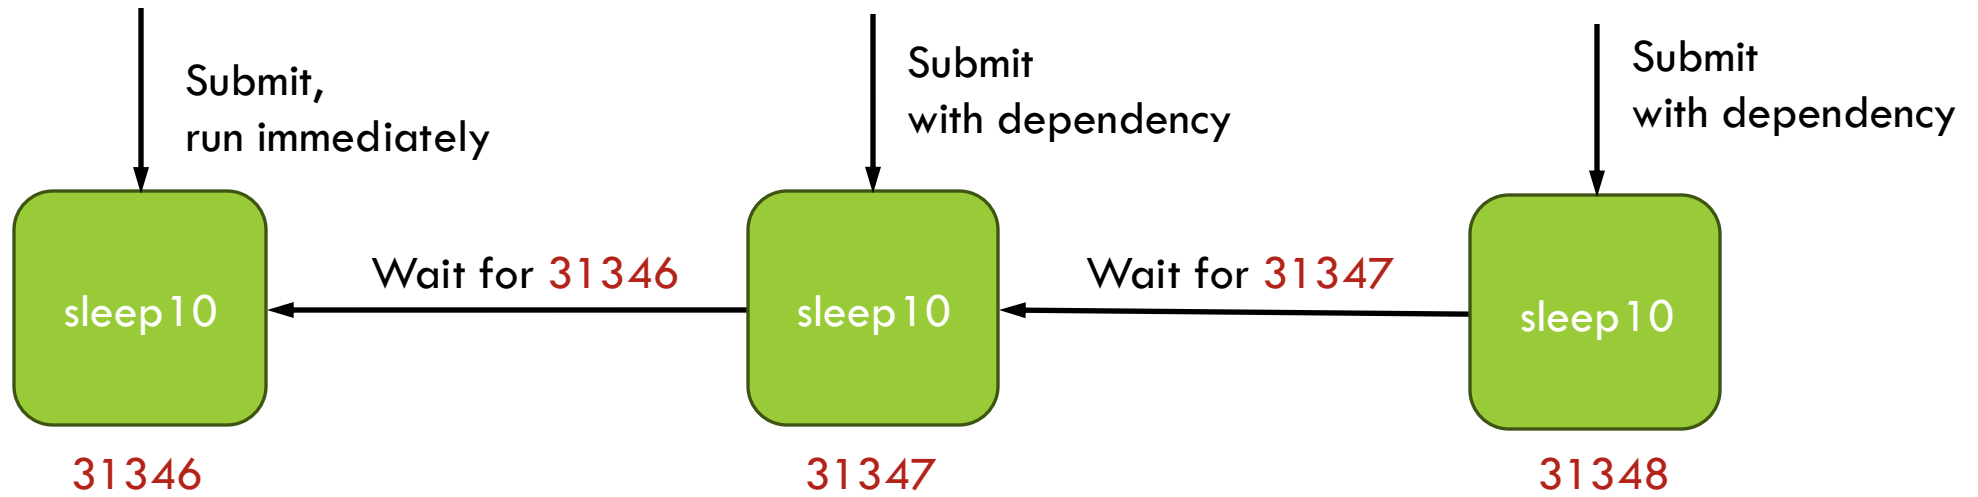
**afterany:job\_id** – Job with job\_id completes for any reason

See sbatch documentation for more complex forms

# SIMPLE PIPELINE WITH DEPENDENCIES

```
$ sbatch sleep10.sh  
Submitted job id 31346  
$ sbatch --dependency=afterok:31346 sleep10.sh  
Submitted job id 31347  
$ sbatch --dependency=afterok:31347 sleep10.sh  
Submitted job id 31348
```

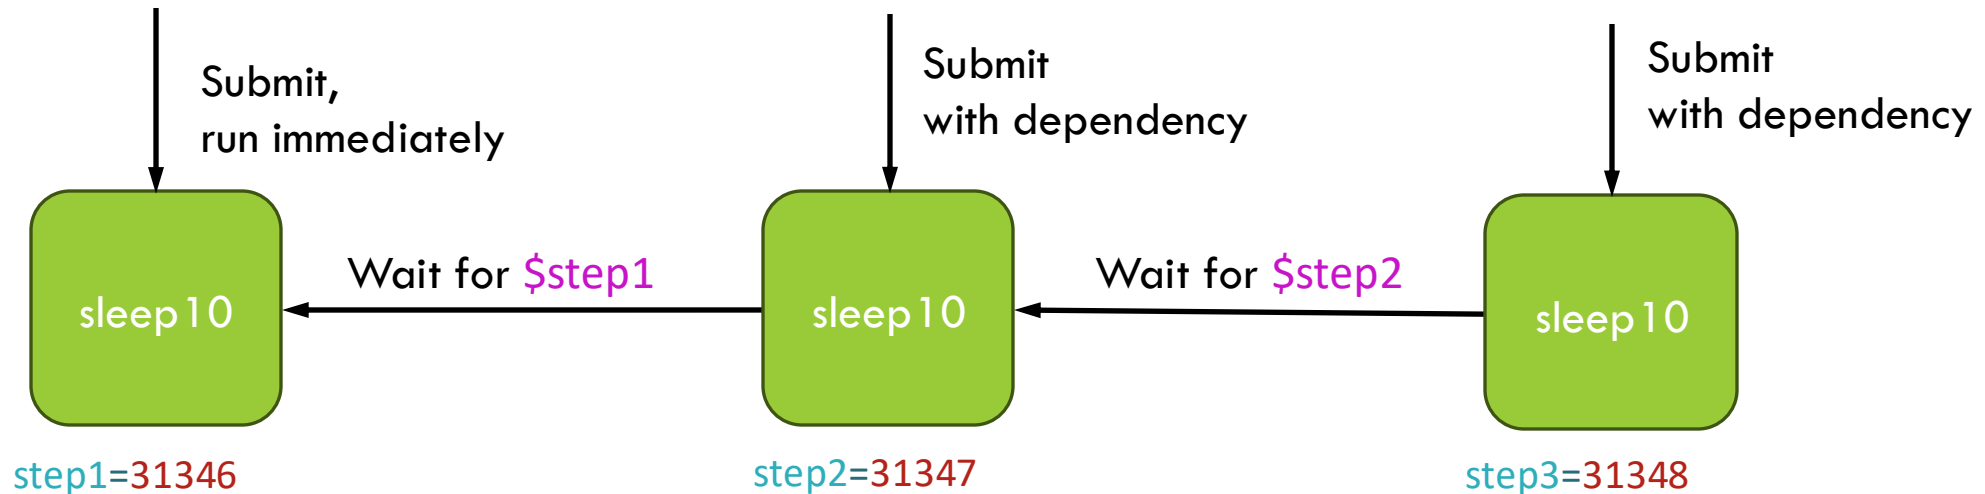




```
(base) [dking@login001 simple1]$ bash simple_pipeline_riviera.sh
```

# SCRIPTABLE PIPELINE

```
step1=$(sbatch --parsable sleep10.sh)  
step2=$(sbatch --parsable --dependency=afterok:$step1 sleep10.sh)  
step3=$(sbatch --parsable --dependency=afterok:$step2 sleep10.sh)
```



# ARRAY WITH CONVERGENCE

## Array Script

Wait random number of seconds, print out message

```
1 echo "slept $rand seconds" → $SLURM_ARRAY_JOB_ID/$SLURM_ARRAY_TASK_ID.txt
2 echo "slept $rand seconds" → $SLURM_ARRAY_JOB_ID/$SLURM_ARRAY_TASK_ID.txt
3 echo "slept $rand seconds" → $SLURM_ARRAY_JOB_ID/$SLURM_ARRAY_TASK_ID.txt
4 echo "slept $rand seconds" → $SLURM_ARRAY_JOB_ID/$SLURM_ARRAY_TASK_ID.txt
```

## Convergence script

Waits for all array jobs to finish

```
cat $SLURM_ARRAY_JOB_ID/*.txt
```

slept 15 seconds  
slept 11 seconds  
slept 16 seconds  
slept 14 seconds

Wait random number of seconds, print out message

```
1 echo "slept $rand seconds" —→ $SLURM_ARRAY_JOB_ID/$SLURM_ARRAY_TASK_ID.txt
2 echo "slept $rand seconds" —→ $SLURM_ARRAY_JOB_ID/$SLURM_ARRAY_TASK_ID.txt
3 echo "slept $rand seconds" —→ $SLURM_ARRAY_JOB_ID/$SLURM_ARRAY_TASK_ID.txt
4 echo "slept $rand seconds" —→ $SLURM_ARRAY_JOB_ID/$SLURM_ARRAY_TASK_ID.txt
```

Waits for all array jobs to finish

```
cat $SLURM_ARRAY_JOB_ID/*.txt
```

```
job_monitor.sh 31588,31589
```

```
queue -u dking -j 31588,31589
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	ODELIST(REASON)
31588_[1-10]	short-cpu	array	dking	PD	0:00	1	(None)
31589	short-cpu	converge	dking	PD	0:00	1	(Dependency)

```
sacct -X -j 31588,31589
```

JobID	JobName	AllocCPU	State	Elapsed	Timelimit	Start	End	Reason
-------	---------	----------	-------	---------	-----------	-------	-----	--------

CTRL-C to quit.


# ARRAY/CONVERGENCE LAUNCHER SCRIPT

```
#!/usr/bin/env bash
```

```
array_job_id=$(sbatch --parsable --array=1-10 array.sbatch)
```

```
converge_job_id=$(sbatch --parsable --dependency=afterok:${array_job_id} converge.sbatch $array_job_id)
```

\*This script is run with `bash` not `sbatch`



# EXAMPLES: LOOPS TO ARRAYS

Some examples of how to modify your scripts to go from sequential to array



# CHANGING YOUR SCRIPT TO AN ARRAY SCRIPT

Three ways to modify/remove loops

1. Arguments on the command line
2. Already looping over an array
3. Reading a file, use sed to select specific line

# CHANGING YOUR SCRIPT TO AN ARRAY SCRIPT

## 1. Arguments on the command line


Original: script loops over '\$@'

```
for arg in "$@"; do
    progname $arg
...
done
```

New version:

- convert command line to array
- Set loop variable from array and `$SLURM_ARRAY_TASK_ID`
- Remove loop

```
args=($@)
arg=${args[$SLURM_ARRAY_TASK_ID]}
#for arg in "$@"; do
    progname $arg
...
#done
```



```
args=($@)
arg=${args[$SLURM_ARRAY_TASK_ID]}
progname $arg
...
```

# CHANGING YOUR SCRIPT TO AN ARRAY SCRIPT

## 2. Already looping over an array in main script

Original: '@' expands to full array

```
# Loop through each subdirectory
for folder_name in "${subdirectories[@]}"; do
    # Verify if the folder contains exactly two files: .dv and its reference file
    dv_files=( $(find "$folder_name" -maxdepth 1 -type f -name "*.dv" | sort) )
    ...
    ...
    ...
done
```

New version: Keep the loop, replace with `$SLURM_ARRAY_TASK_ID`

```
# Loop through each subdirectory
for folder_name in "${subdirectories[$SLURM_ARRAY_TASK_ID]}"; do
    # Verify if the folder contains exactly two files: .dv and its reference file
    dv_files=( $(find "$folder_name" -maxdepth 1 -type f -name "*.dv" | sort) )
    ...
    ...
    ...
done
```

# CHANGING YOUR SCRIPT TO AN ARRAY SCRIPT


## 3. Select a line of a file for input with **sed**

Original: while loop through file

```
while read -a line; do
  arg1=${line[0]}
  arg2=${line[1]}
  arg3=${line[2]}
  cmd $arg1 $arg2 $arg3
done < metadata.txt
```

New version: Pass **\$SLURM\_ARRAY\_TASK\_ID** to **sed**

```
line=( $(sed -n "${SLURM_ARRAY_TASK_ID}p" metadata.txt) )
arg1=${line[0]}
arg2=${line[1]}
arg3=${line[2]}
cmd $arg1 $arg2 $arg3
```



# DEBUGGING WITH EXIT TRAP

Cool stuff

# DEBUGGING/COOL STUFF

script

```
#!/usr/bin/env bash
```

```
echo "<job started on $(date)>"
```

cmd1

cmd2

cmd3

cmd4

```
echo "<job finished on $(date)>"
```

Output/log if successful

```
<job started on Fri Dec 6 13:25>
```

```
[ cmd1 output ]
```

```
[ cmd2 output ]
```

```
[ cmd3 output ]
```

```
[ cmd4 output ]
```

```
<job finished on Fri Dec 6 15:25>
```

An error early on

```
<job started on Fri Dec 6 13:25>
```

```
[ cmd1 output ]
```

```
cmd2 error output cmd2 error cmd2 error output  
cmd2 error output cmd2 error output  
cmd3 error output cmd3 error cmd3 error output  
cmd3 error output cmd3 error output  
cmd4 error output cmd4 error cmd4 error output  
cmd4 error output cmd4 error output
```

# DEBUGGING/COOL STUFF — USING BASH SETTINGS

script

```
#!/usr/bin/env bash
set -e # quit on error
set -o pipefail # detect error in pipe
echo "<job started on $(date)>"
```

cmd1

cmd2

cmd3

cmd4

```
echo "<job finished on $(date)>"
```

An error early on

<job started on Fri Dec 6 13:25>

[ cmd1 output ]

cmd2 error output cmd2 error cmd2 error  
or output cmd2 error output

# DEBUGGING/COOL STUFF — THE EXIT SIGNAL

script

```
#!/usr/bin/env bash
set -e # quit on error
set -o pipefail # detect error in pipe
echo "<job started on $(date)>"
```

cmd1

cmd2

cmd3

cmd4

```
echo "<job finished on $(date)>"
```

An error early on

<job started on Fri Dec 6 13:25>

[ cmd1 output ]

cmd2 error output cmd2 error cmd2 error  
or output cmd2 error output

**<job FAILED on Fri Dec 6 13:26>**



# DEBUGGING/COOL STUFF — TRAPPING SIGNALS

Signals can be sent to any process on linux to interrupt, cancel jobs, among other things

Some signals can be handled “trapped” and given code to run

## **Syntax:**

```
trap code_to_run SIGNAL
```

## **Example:**

Script says “have a nice day” when user CTRL-C’s their script:

```
trap 'echo "have a nice day"; exit 0' SIGINT
```

# DEBUGGING/COOL STUFF — THE EXIT TRAP

script

```
#!/usr/bin/env bash
set -e # quit on error
set -o pipefail # detect error in pipe

echo "<job started on $(date)>"
trap 'echo "Job ended at $(date) "' EXIT

cmd1

cmd2

cmd3

cmd4

echo "<job finished on $(date)>"
```

An error early on

<job started on Fri Dec 6 13:25>

[ cmd1 output ]

cmd2 error output cmd2 error cmd2 error  
or output cmd2 error output

<job ended at Fri Dec 6 13:26>

# DEBUGGING/COOL STUFF — THE EXIT TRAP

script

```
#!/usr/bin/env bash
set -e # quit on error
set -o pipefail # detect error in pipe

echo "<job started on $(date)>"
trap 'echo "Job ended at $(date) "' EXIT

cmd1

cmd2

cmd3

cmd4

echo "<job finished on $(date)>"
```

Output/log if successful

```
<job started on Fri Dec 6 13:25>

[ cmd1 output ]

[ cmd2 output ]

[ cmd3 output ]

[ cmd4 output ]

<job finished on Fri Dec 6 15:25>
<job ended at Fri Dec 6 15:25>
```

# EXIT TRAP WITH HANDLING FUNCTION

```
#!/usr/bin/env bash
echo "##### Script started at $(date) >>>>>>>"
set -e # exit on error
JOBSTEP=SETUP

# error handling
davidExitFunc()
{
    exitcode=$1
    if [ -z "$exitcode" ] || [ $exitcode -eq 0 ] # wasn't provided or is 0
    then
        echo "##### Script reached $JOBSTEP successfully at $(date) >>>>>>>"
    else
        echo "##### Script failed at $JOBSTEP with exit code $exitcode at $(date) >>>>>>>"
    fi
}

trap 'davidExitFunc $?' EXIT

sleep 1
# BEGIN WORKFLOW
JOBSTEP=STEP1
# do something
JOBSTEP=STEP2
[ 1 -lt 4 ] # delete a number to trigger a syntax error

JOBSTEP=STEP3
# exit 1 # uncomment to exit in controlled manner

JOBSTEP=STEP4
# stop prematurely but without error, perhaps debugging
exit 0

JOBSTEP=END
```

**# quits at exit 1 under JOBSTEP=STEP3**

\$ bash cleanup\_function.sh

##### Script started at Mon Dec 9 11:51 >>>>>>>

##### Script failed at STEP3 with exit code 1 at Mon Dec 9 11:51:49  
>>>>>>>

**# delete the '4' under JOBSTEP=STEP2 (syntax error)**

\$ bash cleanup\_function.sh

##### Script started at Mon Dec 9 11:53 >>>>>>>

cleanup\_function.sh: line 25: [: 1: unary operator expected  
##### Script failed at STEP2 with exit code 2 at Mon Dec 9 11:53:42  
>>>>>>>

**# replace the 4 and comment out exit statements**

\$ bash cleanup\_function.sh

##### Script started at Mon Dec 9 11:58:11 >>>>>>>

##### Script reached END successfully at Mon Dec 9 11:58:12 >>>>>>>

# EXIT TRAP WITH HANDLING FUNCTION

## Other uses:

- Close connections, free resources
- Ex: Close Amazon machine instance (charges MONEY \$\$\$)

```
#!/bin/bash  
  
# define the base AMI ID somehow  
ami=$1  
  
function finish {  
    ec2-terminate-instances "$ami"  
}  
  
trap finish EXIT  
  
ec2-run-instances "$ami"
```

Simplified from: <http://redsymbol.net/articles/bash-exit-traps/>

# REPO FOR THIS TALK AND SLURM TOOLS

<https://github.com/Colorado-State-University-CMB/slurm-scripting-pipelines>

- Use: `git clone https://github.com/Colorado-State-University-CMB/slurm-scripting-pipelines` on alpine or rivia
- Do `bash INSTALL.sh` to install `job_monitor.sh`, which is used in the pipeline examples
- The subdirectories contain the examples used in this presentation